

Deep Learning for Computer Vision

Image Sampling and Interpolation

Vineeth N Balasubramanian

Department of Computer Science and Engineering
Indian Institute of Technology, Hyderabad



Cost Improvement using Convolution Theorem?

Convolution Theorem

- Fourier transform of convolution of two functions is product of their Fourier transforms:

$$F[g * h] = F[g]F[h]$$

- Convolution** in spatial domain can be obtained through **multiplication** in frequency domain!

$$g * h = F^{-1}[F[g]F[h]]$$

NPTEL

Cost Improvement using Convolution Theorem?

Convolution Theorem

- Fourier transform of convolution of two functions is product of their Fourier transforms:

$$F[g * h] = F[g]F[h]$$

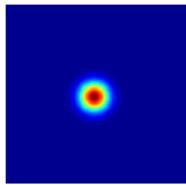
- Convolution** in spatial domain can be obtained through **multiplication** in frequency domain!

$$g * h = F^{-1}[F[g]F[h]]$$

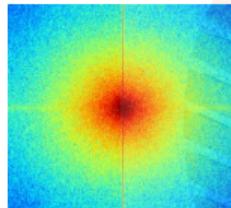
- Image convolution needs $O(N^2 \cdot k^2)$ time, where $N \times N$ is image size, and $k \times k$ is kernel size
- By performing convolution in Fourier domain, cost is: $O(N^2)$ for a single pass over the image + cost of FFT: $O(N^2 \log N^2)$ for the image and $O(k^2 \log k^2)$ for the kernel $\approx O(N^2 \log N^2 + k^2 \log k^2)$, in total (other terms additive)

Exercise: Match spatial domain image to Fourier magnitude image

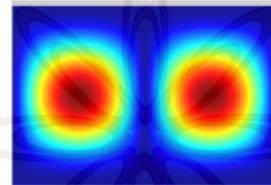
1



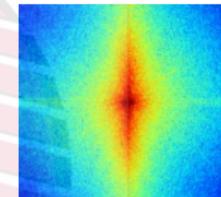
2



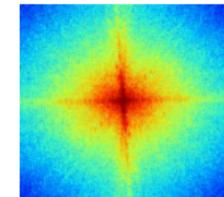
3



4



5



B



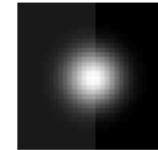
C



A



D

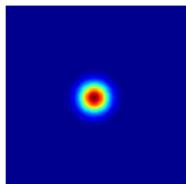


E

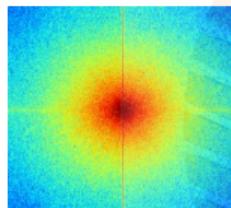


Exercise: Match spatial domain image to Fourier magnitude image

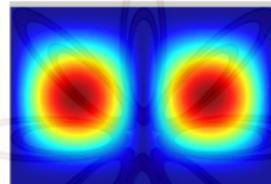
1



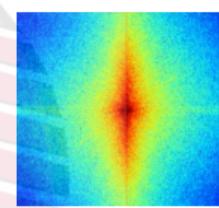
2



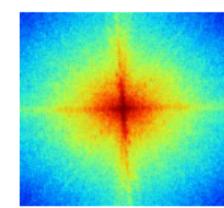
3



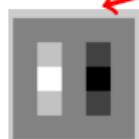
4



5



A



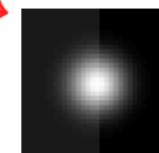
B



C



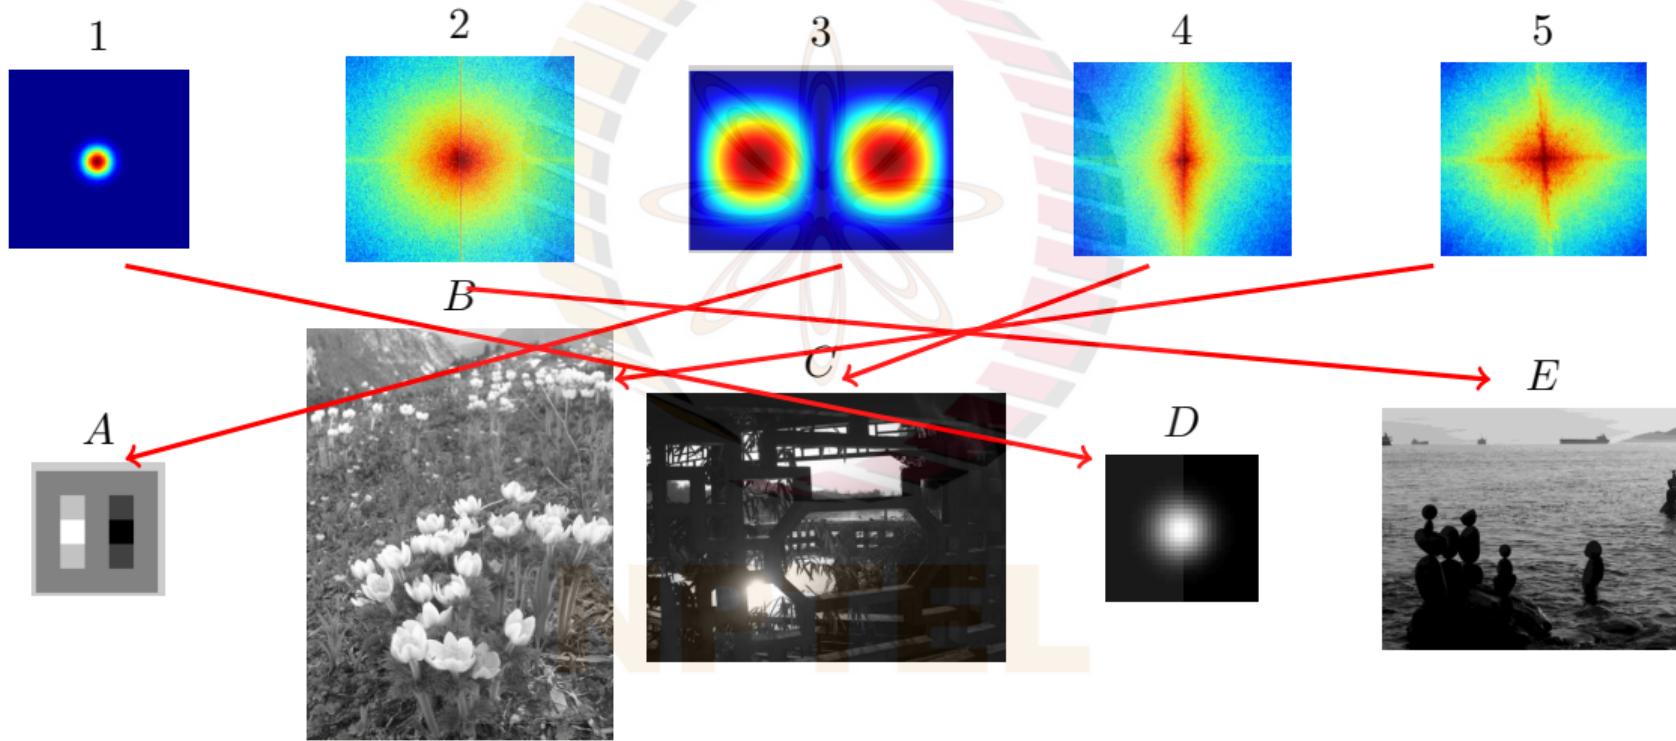
D



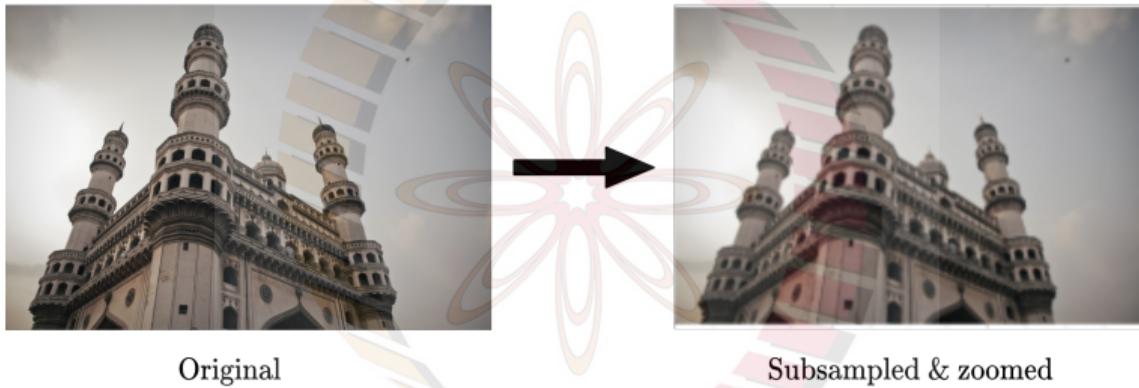
E



Exercise: Match spatial domain image to Fourier magnitude image



What sense does a low-resolution image make to us?

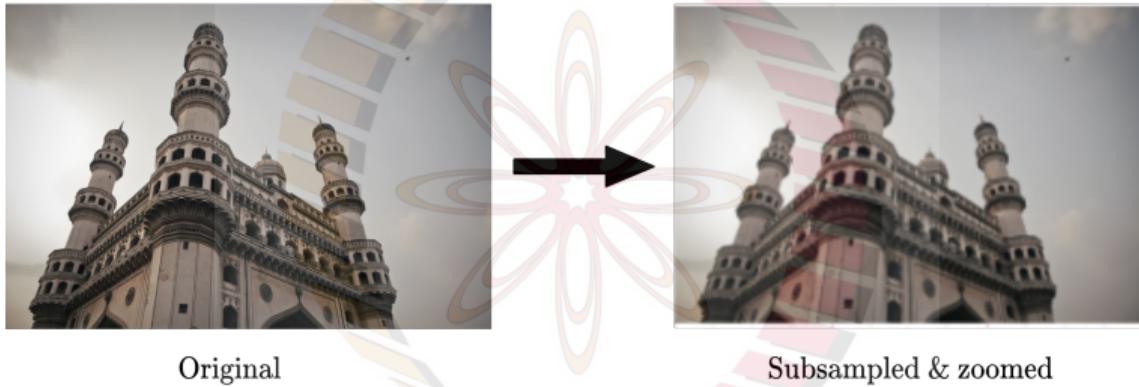


Clues from human perception:

- Early processing in human's filters for various orientations and scales of frequency.
- Perceptual cues in mid-high frequencies dominate perception.
- When we see an image from far away, we are effectively **sub-sampling** it.

Credit: Ron Hansen (Unsplash)

What sense does a low-resolution image make to us?



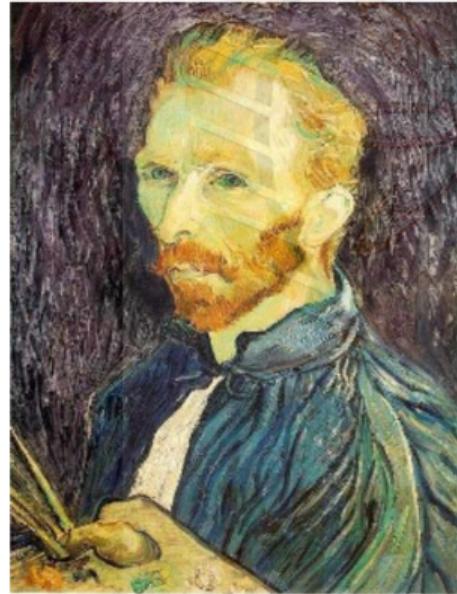
Clues from human perception:

- Early processing in human's filters for various orientations and scales of frequency.
- Perceptual cues in mid-high frequencies dominate perception.
- When we see an image from far away, we are effectively **sub-sampling** it.

Credit: Ron Hansen (Unsplash)

Sub-sampling

Throw away every other row and column to create a 1/2 size image.



1/2



1/4

Credit: S Seitz, R Urtasun

Sub-sampling

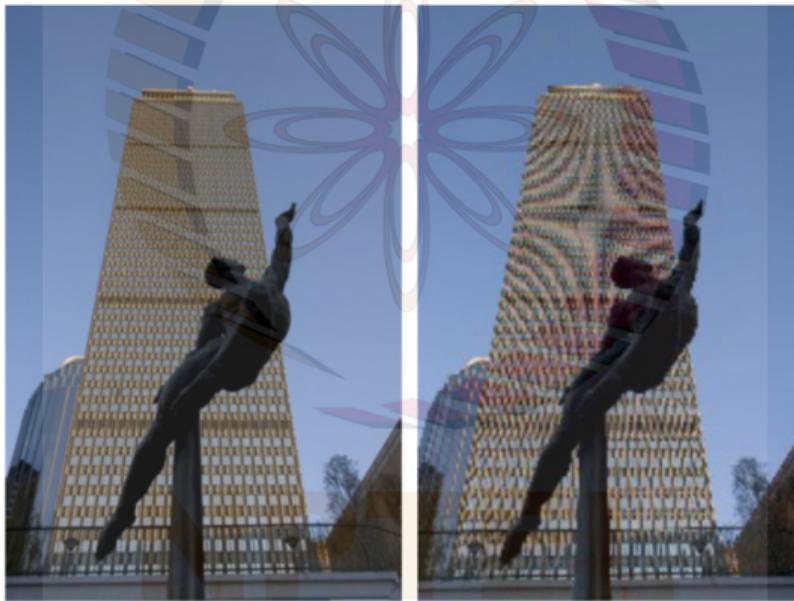
Why does this look so crusty?



Credit: S Seitz, R Urtasun

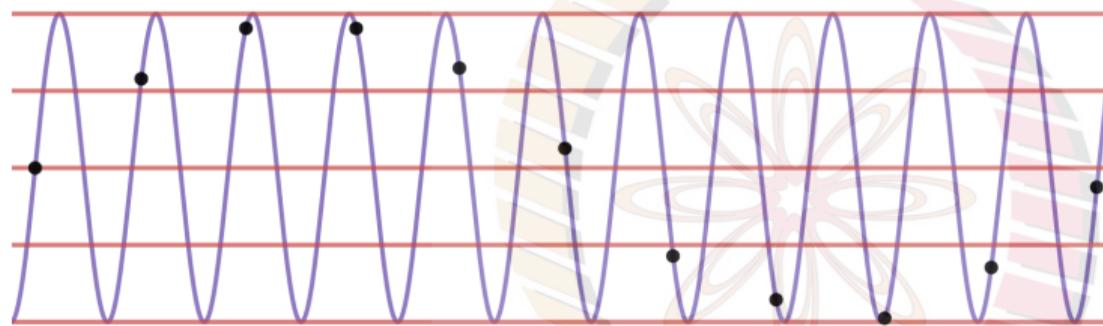
Sub-sampling

What's happening?



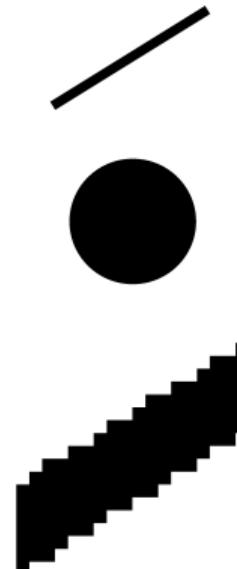
Credit: S Seitz, R Urtasun

Aliasing

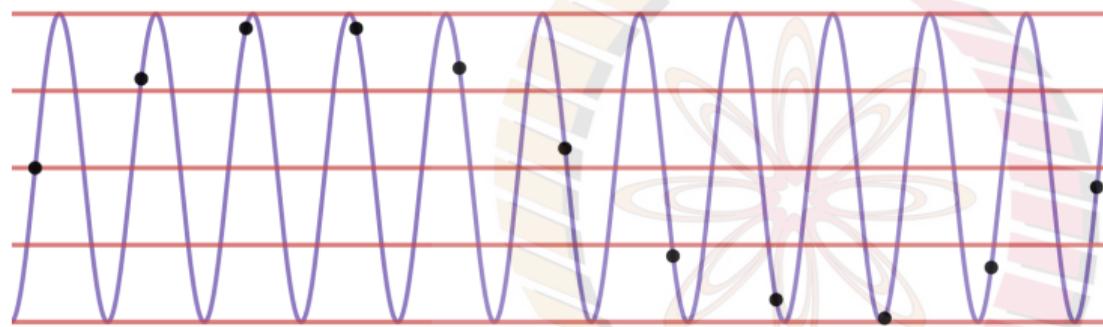


Aliased

- Occurs when your sampling rate is not high enough to capture the amount of detail in your image.
- To do sampling right, need to understand the structure of your signal/image.
- The minimum sampling rate is called the Nyquist rate.

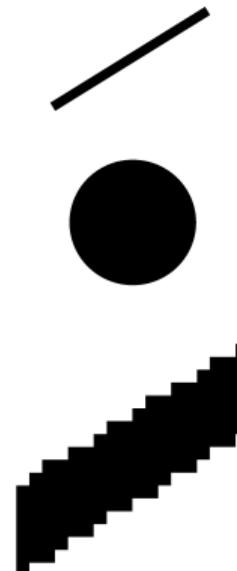


Aliasing

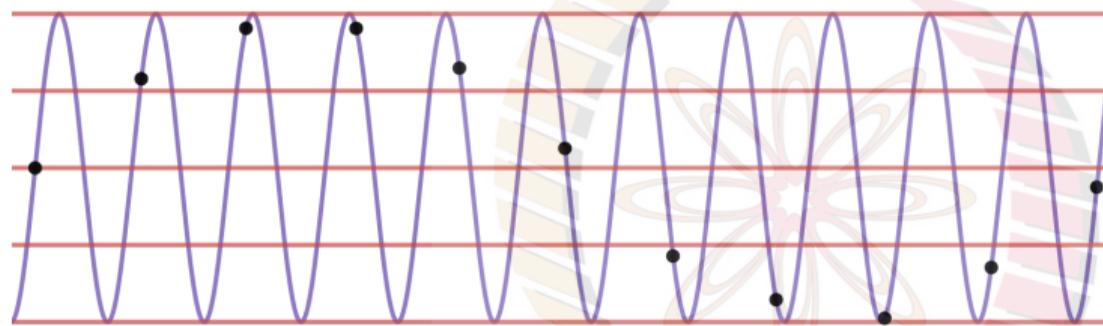


Aliased

- Occurs when your sampling rate is not high enough to capture the amount of detail in your image.
- To do sampling right, need to understand the structure of your signal/image.
- The minimum sampling rate is called the Nyquist rate.

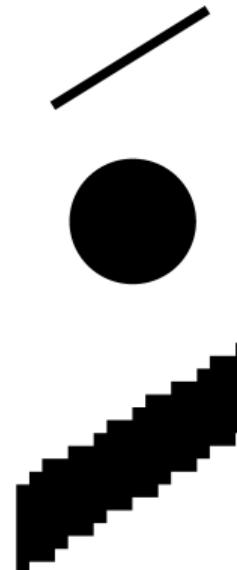


Aliasing



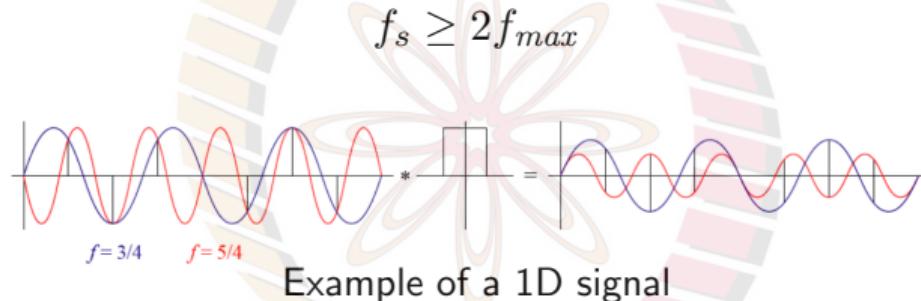
Aliased

- Occurs when your sampling rate is not high enough to capture the amount of detail in your image.
- To do sampling right, need to understand the structure of your signal/image.
- The minimum sampling rate is called the **Nyquist rate**.



Aliasing: Problems

Shannon's Sampling Theorem shows that the minimum sampling is:

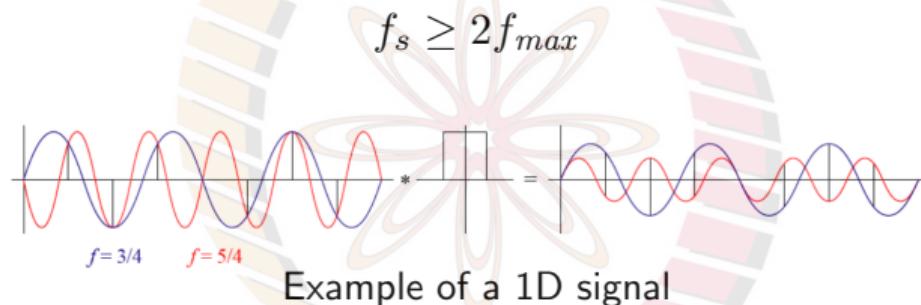


Examples

- **Image**
 - Striped shirt's pattern look weird on screen.
- **Video**
 - **Wagon Wheel effect:** Wheels spins in the opposite direction at high speed.
- **Graphics**
 - Checkerboards disintegrate in ray tracing.

Aliasing: Problems

Shannon's Sampling Theorem shows that the minimum sampling is:

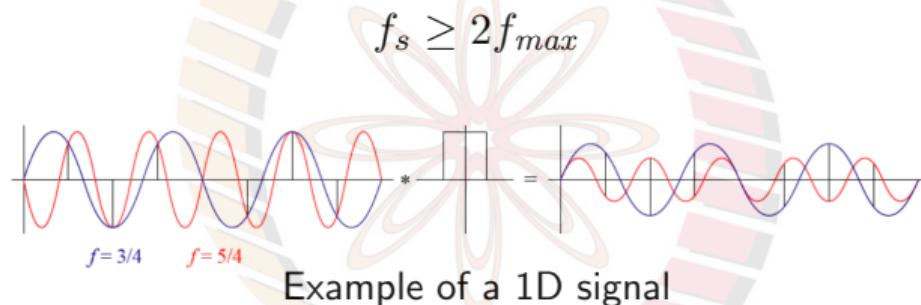


Examples

- **Image**
 - Striped shirt's pattern look weird on screen.
- **Video**
 - **Wagon Wheel effect:** Wheels spins in the opposite direction at high speed.
- **Graphics**
 - Checkerboards disintegrate in ray tracing.

Aliasing: Problems

Shannon's Sampling Theorem shows that the minimum sampling is:



Examples

- **Image**

- Striped shirt's pattern look weird on screen.

- **Video**

- **Wagon Wheel effect:** Wheels spins in the opposite direction at high speed.

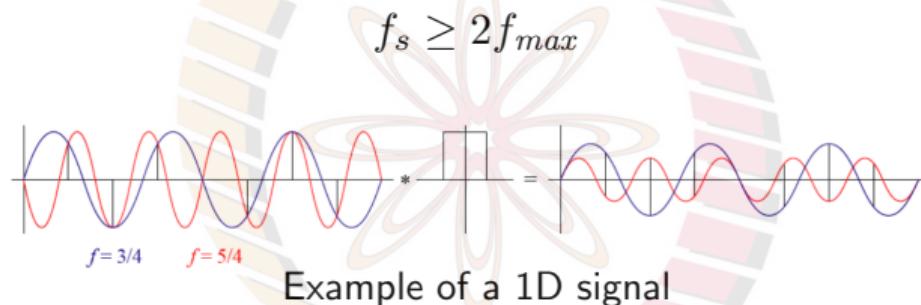
- **Graphics**

- Checkerboards disintegrate in ray tracing.

NPTEL

Aliasing: Problems

Shannon's Sampling Theorem shows that the minimum sampling is:



Examples

- **Image**
 - Striped shirt's pattern look weird on screen.
- **Video**
 - **Wagon Wheel effect:** Wheels spins in the opposite direction at high speed.
- **Graphics**
 - Checkerboards disintegrate in ray tracing.

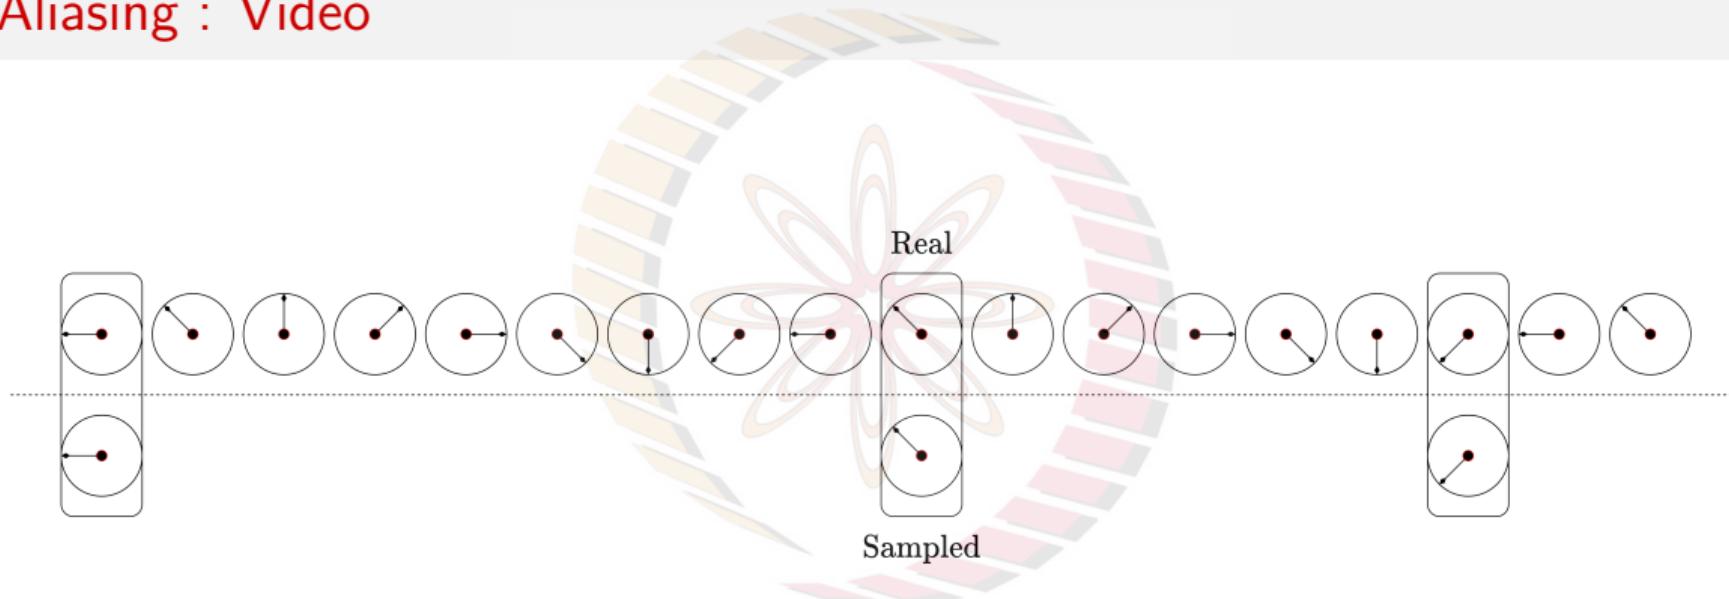


Aliasing: Image



Striped shirt's pattern look weird on screen.

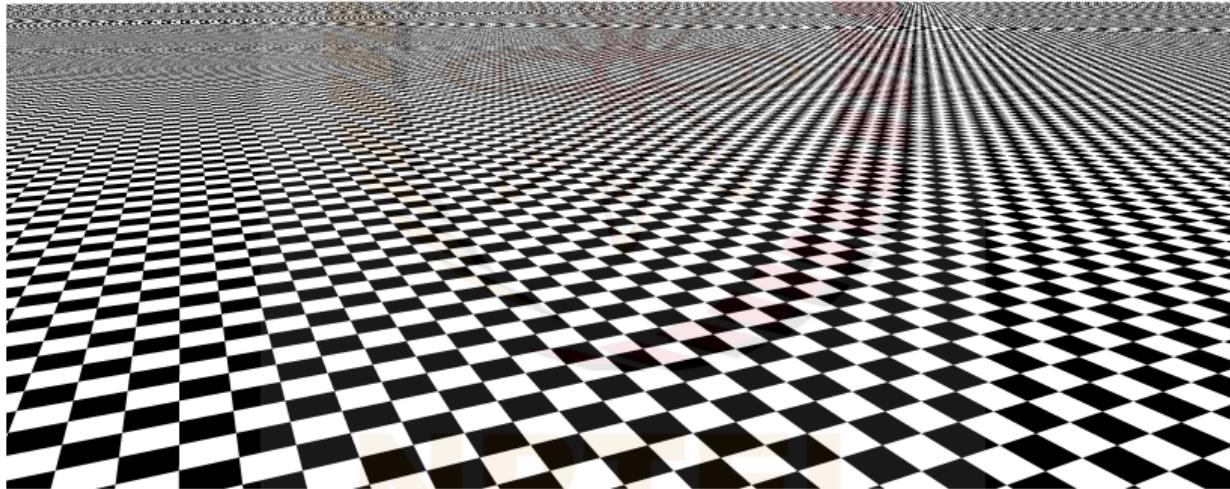
Aliasing : Video



Wagon Wheel effect: Wheels spins in the opposite direction at high speed.

NPTEL

Aliasing: Graphics



Checkerboards disintegrate in ray tracing.

Aliasing: Nyquist Limit 2D example



Credit: S Seitz, R Urtasun

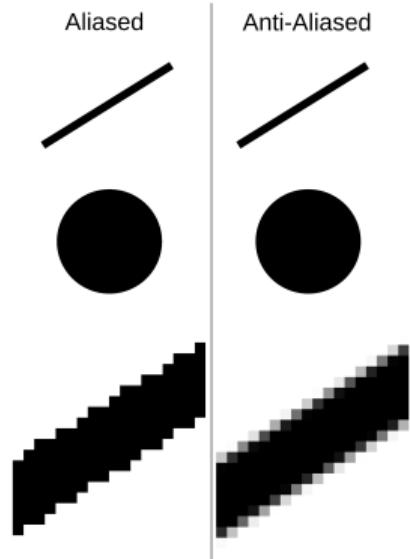
Anti-aliasing



NPTEL

Anti-aliasing

Example: Gaussian Pre-filtering



Credit: N Snavely, R Urtasun

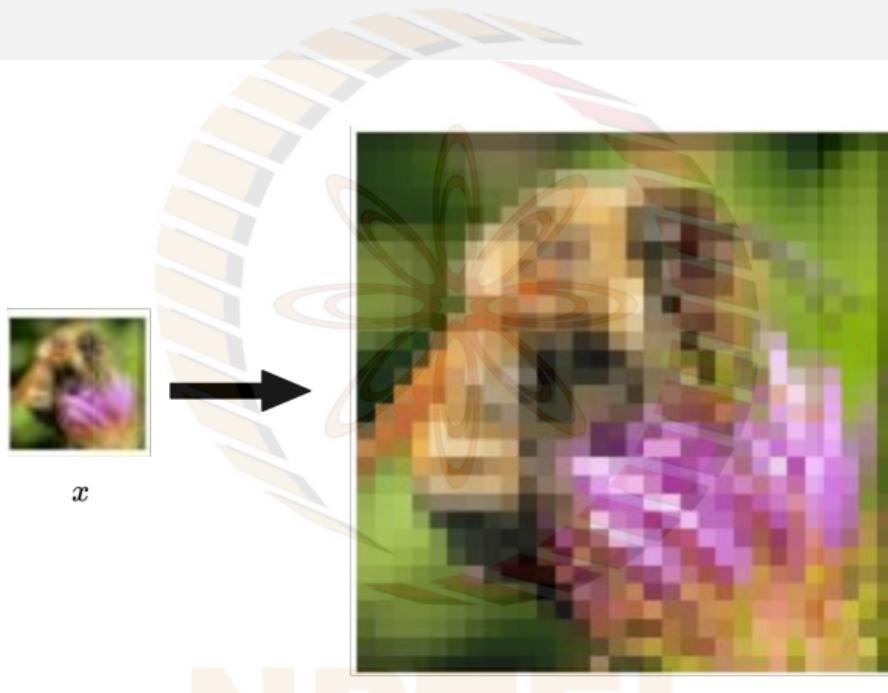


Subsampling with Gaussian Pre-filtering



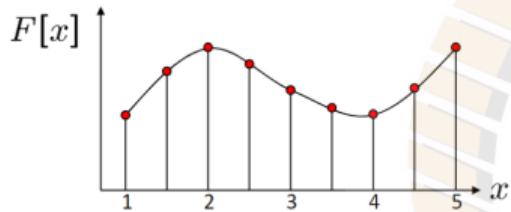
Credit: N Snavely, R Urtasun

Upsampling



How to go from left to right? **Interpolation**. Simple method: Repeat each row and column 10 times (Nearest Neighbour Interpolation).

Interpolation

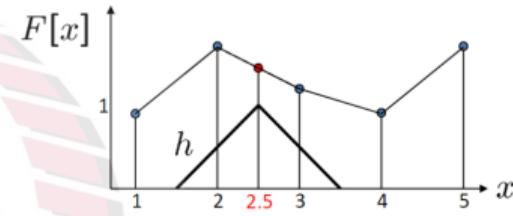


Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.

NPTEL



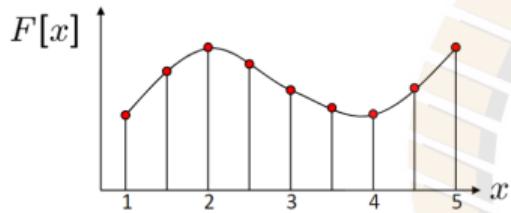
What if we don't know f ?

- Guess an approximation: Can be done in a principled way via filtering.
- Convert F to a continuous function:

$$f_F(x) = \begin{cases} F\left(\frac{x}{d}\right) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct: $\hat{f} = h * f_F$

Interpolation

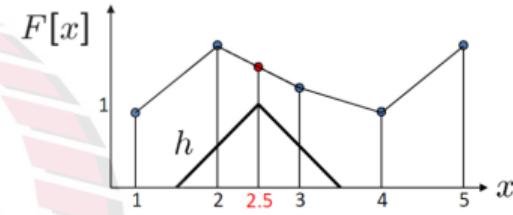


Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.

NPTEL



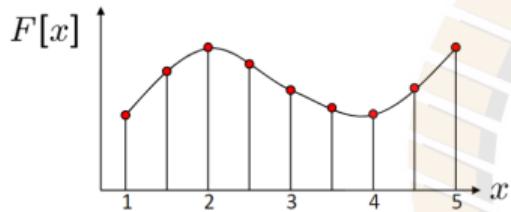
What if we don't know f ?

- Guess an approximation: Can be done in a principled way via filtering.
- Convert F to a continuous function:

$$f_F(x) = \begin{cases} F\left(\frac{x}{d}\right) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct: $\hat{f} = h * f_F$

Interpolation

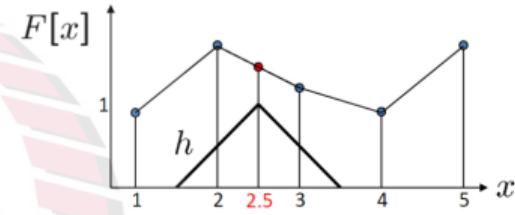


Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.

IITTEL



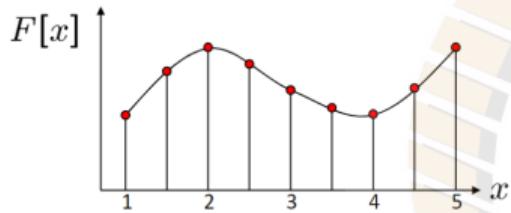
What if we don't know f ?

- Guess an approximation: Can be done in a principled way via filtering.
- Convert F to a continuous function:

$$f_F(x) = \begin{cases} F\left(\frac{x}{d}\right) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct: $\hat{f} = h * f_F$

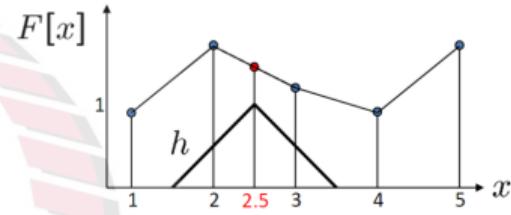
Interpolation



Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.



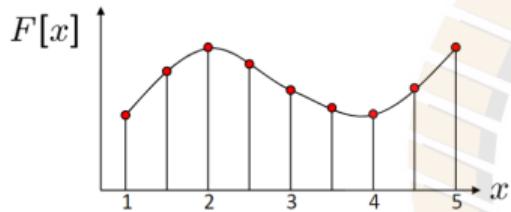
What if we don't know f ?

- Guess an approximation: Can be done in a principled way via filtering.
- Convert F to a continuous function:

$$f_F(x) = \begin{cases} F\left(\frac{x}{d}\right) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct: $\hat{f} = h * f_F$

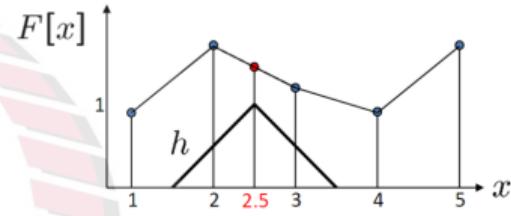
Interpolation



Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.



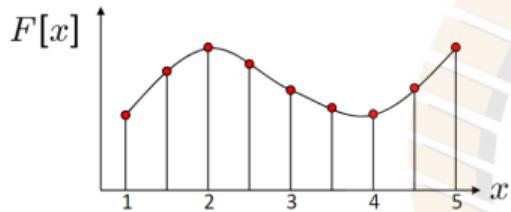
What if we don't know f ?

- Guess an approximation: Can be done in a principled way via filtering.
- Convert F to a continuous function:

$$f_F(x) = \begin{cases} F\left(\frac{x}{d}\right) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct: $\hat{f} = h * f_F$

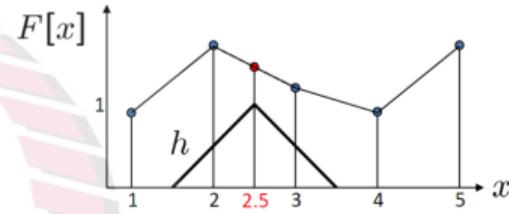
Interpolation



Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.



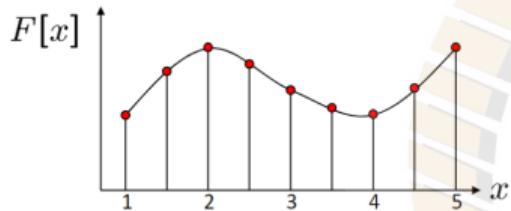
What if we don't know f ?

- Guess an approximation: Can be done in a principled way via filtering.
- Convert F to a continuous function:

$$f_F(x) = \begin{cases} F\left(\frac{x}{d}\right) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct: $\hat{f} = h * f_F$

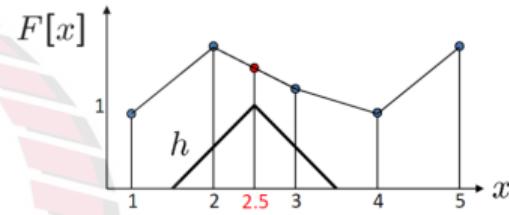
Interpolation



Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.



What if we don't know f ?

- Guess an approximation: Can be done in a principled way via filtering.
- Convert F to a continuous function:

$$f_F(x) = \begin{cases} F\left(\frac{x}{d}\right) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct: $\hat{f} = h * f_F$

Interpolation as Convolution

- To **interpolate** (or upsample) an image to a higher resolution, we need an **interpolation kernel** with which to **convolve** the image:

$$g(i, j) = \sum_{k,l} f(k, l)h(i - rk, j - rl)$$

Above formula similar to discrete convolution^a, except that we replace k and l in $h(\cdot)$ with rk and rl , where r is the upsampling rate.

- Linear interpolator (corresponding to *tent kernel*) produces interpolating piecewise linear curves.
- More complex kernels e.g., B-splines.

^a $g = f * h \implies g(i, j) = \sum_{k,l} f(k, l)h(i - k, j - l)$

Interpolation as Convolution

- To **interpolate** (or upsample) an image to a higher resolution, we need an **interpolation kernel** with which to **convolve** the image:

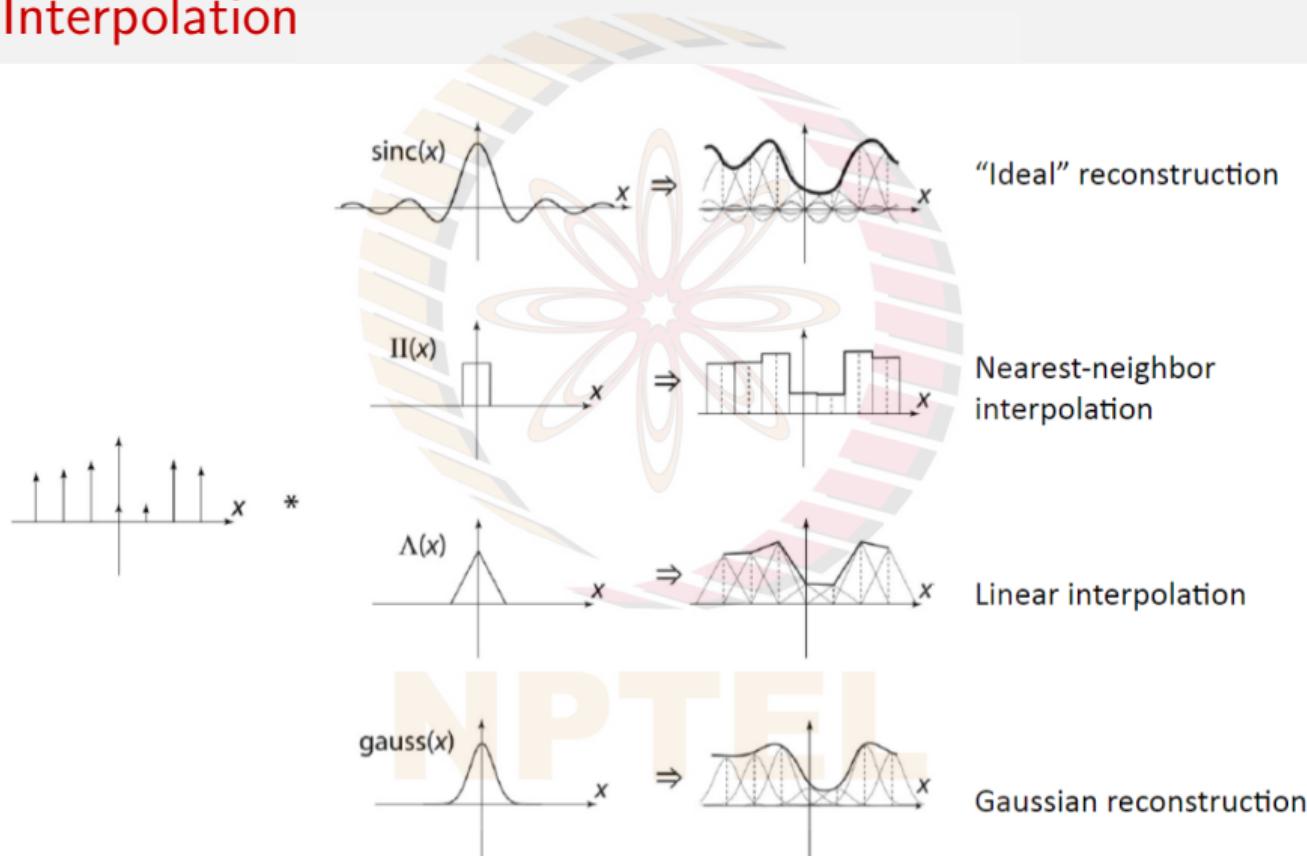
$$g(i, j) = \sum_{k,l} f(k, l)h(i - rk, j - rl)$$

Above formula similar to discrete convolution^a, except that we replace k and l in $h(\cdot)$ with rk and rl , where r is the upsampling rate.

- Linear interpolator (corresponding to *tent kernel*) produces interpolating piecewise linear curves.
- More complex kernels e.g., B-splines.

^a $g = f * h \implies g(i, j) = \sum_{k,l} f(k, l)h(i - k, j - l)$

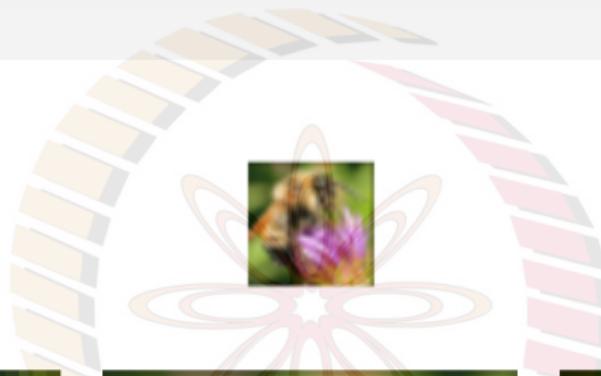
Types of Interpolation



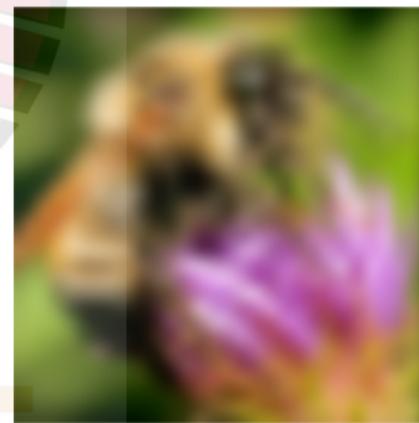
Credit: B Curless

Examples

Original Image:



Upsampled Images:



Left to right: Nearest Neighbour Interpolation, Bilinear Interpolation, Bicubic Interpolation.

Interpolation and Decimation

Interpolation

To **interpolate** (or upsample) an image to a higher resolution, we need an **interpolation kernel** with which to convolve the image (r is upsampling rate):

$$g(i, j) = \sum_{k,l} f(k, l) h(i - rk, j - rl)$$

Decimation (Sub-sampling)

To **decimate** (or sub-sample) an image to a lower resolution, we need an **decimation kernel** with which to convolve the image (r is downsampling rate):

$$g(i, j) = \sum_{k,l} f(k, l) h\left(i - \frac{k}{r}, j - \frac{l}{r}\right)$$

Homework

Readings

- Chapter 3 (§3.5.1-3.5.2), Szeliski, *Computer Vision: Algorithms and Applications*
- Chapter 7 (§7.4), Forsyth and Ponce, *Computer Vision: A Modern Approach*

The NPTEL logo consists of the letters "NPTEL" in a large, bold, sans-serif font. The letters are colored in a gradient of yellow and orange, with the "N" being yellow and the "PTEL" being orange.